

# COMP3301 Assignment 3

## OpenBSD VHD Kernel Driver - Filing the System

Due: 3pm Monday in Week 13 (21st of October)

Submission: Git

Code submission is marked in your prac session in week 13

Last Updated: October 12, 2024

## 1 Academic Integrity

All assessments are individual. You should feel free to discuss aspects of C programming and assessment specifications with fellow students and discuss the related APIs in general terms. You should not actively help (or seek help from) other students with the actual coding of your assessment. It is cheating to look at another student's code, and it is cheating to allow your code to be seen or shared in printed or electronic form. You should note that all submitted code will be subject to automated checks for plagiarism and collusion. If we detect plagiarism or collusion (outside of the assessment process), disciplinary proceedings will be initiated against you. If you are a staff member, Do not be tempted to copy and paste code on student misconduct in the following link: <https://eecs.uq.edu.au/current-students/conduct>

### 1.1 Use of AI Tools

All assessment tasks evaluate your ability to write code without the aid of generative Artificial Intelligence. You are advised that the use of AI technologies to assist in the completion of the assignment is **strictly prohibited** and may constitute student misconduct.

## 2 Background

This assignment extends the concepts of VHD disk images as block devices. This is similar to the `vnd(4)` driver, which supports using files containing a disk image as a block device. The purpose of this assignment is to exercise concepts of low-level disk operations and caching in an operating system kernel environment.

From a high-level point of view, a physical disk device presents a sequence of bytes that can be written to or read from, with the ability to quickly seek an arbitrary position and read and write at that point. Note that this is a simplification that ignores that disks address and provide access to blocks of bytes, not individual bytes.

A file on most operating systems offers similar features, i.e., a sequence of bytes that can be accessed by address. Because of these similarities, it is possible for an operating system to provide a common set of operations on both files and disks (e.g., open, close, read, write, seek, etc.) and allow them to be used interchangeably. For example, you could use `tar` to write an archive to a file in a filesystem, or directly to a disk device. `dd`, `cp`, `cat`, etc can read the bytes

from a raw disk into a file or visa versa. However, operating systems generally provide extra functionality on top of disk devices such as the ability to partition disks and mount filesystems from them.

## 2.1 vnd(4)

The `vnd(4)` driver in OpenBSD provides a "disk-like interface to a file". This means the OpenBSD kernel can open a file and present it as a block device to the rest of the system, which in turn allows for the creation and use of filesystems on these disk images.

The `vnd(4)` driver currently only supports using raw disk images as backing files. There's a one-to-one mapping of data offsets for data in the end disk device and the byte offset of that data in the underlying file. This makes the implementation very simple, with the downside that the backing file has to be the same size as the disk `vnd` is presenting. If you have a 32G disk image, the file will be 32G regardless of how much data is actually stored inside a filesystem mounted on top of it. Similar functionality exists in the `loop` driver in Linux, and the `lofi` driver in Solaris and Illu

## 2.2 Virtual Hard Disk

Virtual Hard Disk (VHD) is a format used by Microsoft's Virtual PC and Oracle VM VirtualBox. A defining feature of VHD is that it allows a disk image on demand to be created without pre-allocate space for the entire image, works in a very

VHD also supports using a separate file. When used in a file, only changes made to the disk are written to the file. Due to the popularity of VHD, it is often used to be distributed as VHD



Microsoft's Virtual PC and Oracle VM VirtualBox. A defining feature of VHD is that it allows a disk image on demand to be created without pre-allocate space for the entire image, works in a very

writing changes to a separate file. When used in a file, only changes made to the disk are written to the file. Due to the popularity of VHD, it is often used to be distributed as VHD

The VHD format used by

<https://stluc.manta.uqcloud.net/comp3301/public/2024/comp3301-vhd-spec.pdf>

The same file can be found on Blackboard. This specification differs from the official Microsoft specification in that it includes annotations for ease of understanding and contains corrections to errors in the official specification. The official VHD format is documented by Microsoft in [Virtual Hard Disk Format Spec\\_10\\_18\\_06.doc](#).

## 3 Instructions

To complete the assignment, you will need to do the following:

1. Download the base code [patch](#)

```
cd ~
ftp https://stluc.manta.uqcloud.net/comp3301/public/2024/comp3301-2024-a3.patch
```

2. Create the a3 branch

```
cd /usr/src
git checkout -b a3 openbsd-7.5
```

3. Apply the base code patch

```
git am < ~/comp3301-2024-a3.patch
```

4. Install the includes

```
cd /usr/src/include
doas make includes
```

5. Build the kernel

```
cd /usr/src/sys
make obj
make config
make -j4
doas make insta
```

6. Reboot

```
doas reboot
```

7. Build and install v

```
cd /usr/src/usr
make obj
make
doas make insta
```

8. Create vhd(4) dev

```
doas cp /usr/src/etc/etc.amd64/MAKEDEV /dev
cd /dev
doas chmod 755 MAKEDEV
doas ./MAKEDEV vhd0
```



## 4 Specifications

You will be extending the OpenBSD kernel to add support for using VHD files as a backend for a `vhd(4)` virtual block device. `vhd(4)` is roughly based on `vnd(4)`. Boilerplate code for the device entry points and command line utility will be provided, but you will be implementing the handling of the file and the VHD file format within the provided kernel driver.

Only a subset of the VHD functionality listed in Microsoft's specification of the file format is required. The following functionality is required:

- Read support
- Write support
- Fixed-size images
- Dynamic images

Differencing images do not need to be supported. In addition to supporting the VHD file format, the kernel should implement the following:

- Deny attaching VHD files which are invalid, corrupted or contain features not supported by your kernel driver.
- Deny detaching VHD files when the disk is open unless the force flag is passed to the `VHDIODETACH` `ioctl`.
- Return the name of the VHD file the device was attached to for the `VHDIODEFNAME` `ioctl`.
- Populate a `struct stat` for the currently attached VHD file for the `VHDIODESTAT` `ioctl`.

#### 4.1 VHD Disk I/O

Reads and writes against the VHD backing file. Write the read and write functions for this assignment.

Caching of VHD structure data blocks should be implemented. For example, if you read a sector from the same block, you should reuse the cache. For reads and writes, for example, if you read a sector and then write to the same block, the cache replacement should result in one read and one write. The cache replacement implementation details are left to you. Some form of caching is implemented.


#### 4.2 ioctl interface

The following `ioctls` should only work on the raw partition of the character device associated with each `vhd` disk. Except for `VHDIODEATTACH`, they should only work when a `vhd` disk is attached to a backing file.

##### VHDIODEATTACH

Specify the VHD file to attach as a block device, and parameters for using the disk in the kernel. The `vhd_attach` struct contains the following fields:

- `vhd_file` - The name of the VHD file to attach a `vhd` disk to.
- `vhd_readonly` - The `vhd` disk can be written to when set to 0, and if the VHD file is read-only, the `vhd` disk should fail to attach. The `vhd` disk should be read-only when set to a non-zero value, and the VHD file should be opened read-only.



Use `ioctl` against the VHD disk image. The marks for this assignment are in the QR code.

allocation table) and performance. For example, if you read a sector from the same block, you should reuse the cache. Same goes for mixed reads and writes, for example, if you read a sector and then write to the same block, the cache replacement should result in one read and one write. The cache replacement implementation details are left to you. Some form of caching is implemented.

## VHDIOCDetach

This `ioctl` requests the block device be detached, and the backing file closed. If the disk is in use, the request should fail with `EBUSY` unless the unsigned int `ioctl` argument is set to a non-zero value. A non zero value requests the forced removal of the block device and close of the backing VHD file.

## VHDIOCFNAME

This `ioctl` requests the name of the VHD file used for the currently attached block device. The name should be the same as what was passed as the filename in the `VHDIOCATtach` `ioctl`.

## VHDIOCSTAT

This `ioctl` is equivalent to an `fstat(2)` system call against the backing file.

# 5 Provided Tools/Files

## 5.1 vhdctl

The patch includes source code for `vhd` devices. It is similar to `src/usr.sbin/vhdctl`.

defined above to control the patch is applied in

## 5.2 vhdtool

`vhdtool` allows you to create and manage VHD files. It also allows you to perform basic operations on raw and dynamic disk headers by running the following commands:

raw and VHD. It also allows you to dump their footers. It may be installed

```
cd ~
ftp https://stluc.manta.uqcloud.net/comp3301/public/2024/rawtest.img
doas tar -xpf vhdtool.tar
```

`tool.tar`

## 5.3 rawtest.img

This is a raw disk of size 10 MiB which contains a filesystem with the following files:

<code>-rw-r--r--</code>	1	root	wheel	10	Oct	4	21:37	comp3301.txt
<code>drwxr-xr-x</code>	2	root	wheel	512	Oct	4	21:39	folder/
<code>-rwxr-xr-x</code>	1	root	wheel	6496	Oct	4	21:40	hello*
<code>-rw-r--r--</code>	1	root	wheel	94	Oct	4	21:40	hello.c
<code>-rw-r--r--</code>	1	root	wheel	13	Oct	4	21:36	hello.txt
<code>-rw-r--r--</code>	1	root	wheel	2739	Oct	4	21:38	test.txt

Download:

<https://stluc.manta.uqcloud.net/comp3301/public/2024/rawtest.img>.